



AI reinvention made real

Accenture Innovation Challenge 2026

Team details

TEAM NAME: NewGenLabs



Raghav Bansal
(Team Leader)

College: IIT BHU

Stream: Mechanical Engineering

Year of graduation: 2027

Describe the problem statement (200 words)


Enterprise AI sounds certain — even when facts are missing.

Generative AI is moving into banking, insurance, customer service and operations, where one wrong answer can create financial loss, compliance violations or bad business decisions. The issue is not only hallucination. LLMs can produce fluent claims without proving where the facts came from. A model may use an outdated policy, miss a customer-specific clause, retrieve the wrong document or guess a missing detail.

For low-impact questions, a wrong answer is inconvenient. For high-impact questions — such as whether a customer can close a ₹20 lakh loan without penalty — an unsupported answer can be expensive and irreversible.

Existing AI stacks can retrieve, monitor and filter responses, but retrieval and model confidence do not establish claim-level proof against authoritative enterprise state.


Problem: confidence is not proof.



Proposed solution (200 words)

**ControlPlane.ai — Evidence before answer.
No proof, no confident claim.**

ControlPlane.ai is a risk-aware evidence-enforcement layer for enterprise AI. For each request, a Risk Router decides how strict verification must be. High-impact workflows use a versioned Fact Contract: a policy-defined checklist of the facts, sources and freshness required before an AI claim is allowed. An Evidence Orchestrator retrieves permitted data from systems of record — databases, signed contracts, policy repositories and APIs — and creates a Verified Fact State with value, source, version and timestamp. ControlPlane then drafts the response as atomic claims and routes each claim to the cheapest reliable verifier: database checks for transactions, calculators for arithmetic, tests for code, rules engines for policy and semantic verification only when deterministic checks are insufficient. Claims become PROVEN, CONTRADICTED, INSUFFICIENT, CONFLICTING or STALE. Only proven claims are verbalized by the LLM, behind a final guard against unsafe or biased wording. Missing evidence never becomes proof, especially for negative claims like “no fee applies.” The output includes an evidence receipt showing what was checked, approved, withheld or escalated. Deterministic checks run in milliseconds and low-risk traffic skips the gate — so enforcement doesn't slow the AI, and every answer carries replayable proof.



Why it matters (200 words)

Evidence enforcement is what lets enterprise AI touch money, health and law.

Banks, insurers and regulated enterprises want AI in lending, claims and compliance — exactly where wrong answers are costliest. ControlPlane gives chief risk and compliance officers a control point: AI that proves, abstains or escalates instead of guessing. Customers get answers they can act on; auditors get a replayable evidence receipt behind every response.

One layer addresses all three AI risks. Performance: claims are verified against systems of record before verbalization, and negative claims need affirmative proof. Cost: a risk router sends most traffic through cheap deterministic checks — milliseconds, not extra model calls. Responsibility: authorization-aware retrieval prevents data leakage, and a final guard blocks unsafe or biased wording.

We would measure what enterprises actually fear: false-allow rate on high-impact claims, abstention precision, added latency, and cost per verified answer. In Round 2 we would prototype the loan-foreclosure flow end-to-end — mock core-banking data, a rules engine and semantic verification — demonstrating proven, blocked and escalated claims live.

Impact: AI that earns the right to answer.

The three risks in the brief — one enforcement layer

Answered point by point

“Confidently wrong”

Verify claims **before** verbalizing · negative claims need proof · fail closed

“Quietly expensive”

Cheapest reliable verifier · risk router skips low-risk traffic · milliseconds, not model calls

“Subtly biased / leaking data”

Authorized retrieval · final safety guard · evidence receipt for every answer

We’d measure: false-allow rate · abstention precision · added p95 latency · cost per verified answer

Round 2 prototype: mock core-banking DB + rules engine + semantic verifier — live proven / blocked / escalated demo

Thank you

